

Design and Analysis of Algorithms (II)

Various Problems

Guoqiang Li

School of Software, Shanghai Jiao Tong University

Efficient Problems, Difficult Problems

Efficient Algorithms

We have developed algorithms for

- Finding shortest paths in graphs,
- Minimum spanning trees in graphs,
- Matchings in bipartite graphs,
- Maximum increasing subsequences,
- Maximum flows in networks,
-

All these algorithms are **efficient**, since their time requirement grows as a **polynomial function** (such as n , n^2 , or n^3) of the size of the input.

Exponential Search Space

In these problems we are searching for a solution (path, tree, matching) from among an **exponential** population of possibilities.

All these problems could in principle be solved in **exponential time** by checking through all candidate solutions, one by one.

An algorithm with running time 2^n , or worse, is useless in practice.

The **efficient algorithms** is to find clever ways to bypass **exhaustive search**, using clues from the input to narrow down the search space.

Are there “**search problems**” in which seeking a solution among an **exponential** chaos, and the fastest algorithms for them are **exponential**?


Minimum Spanning Trees

Build a Network

Suppose you are asked to **network** a collection of computers by linking selected pairs of them.

This translates into a graph problem in which

- nodes are computers,
- undirected edges are potential links, each with a **maintenance cost**.




Build a Network

The goal is to

- pick enough of these edges that the nodes are **connected**,
- the total maintenance cost is **minimum**.

One immediate observation is that the optimal set of edges cannot contain a **cycle**.



A Greedy Approach



Kruskal's algorithm starts with the **empty graph** and then selects edges from E according to the following rule.

Repeatedly add the next lightest edge that doesn't produce a cycle.

Example:

Starting with an empty graph and then attempt to add edges in increasing order of weight

$B - C$; $C - D$; $B - D$; $C - F$; $D - F$; $E - F$; $A - D$; $A - B$; $C - E$; $A - C$



A General Kruskal's Algorithm

$X = \{ \};$

repeat until $|X| = |V| - 1;$

 pick a set $S \subset V$ for which X has no edges between S and $V - S;$

 let $e \in E$ be the minimum-weight edge between S and $V - S;$

$X = X \cup \{e\};$

Remarks

Disjoint set is a backbone data structure in the Kruskal's algorithm.

And **amortized analysis** is adopted as an analysis of the algorithm.

Prim's Algorithm

An alternative is **Prim's** algorithm, where the set of edges X always forms a subtree, and S is chosen to be the set of this tree's vertices.

On each iteration, X grows by one edge. We can equivalently think of S as growing to include the vertex $v \notin S$ of smallest cost :

$$\text{cost}(v) = \min_{u \in S} w(u, v)$$

Quiz

Q: Which data structure do you think the Prim's algorithm adopts?

A Little Change of the MST

What if the tree is not allowed to **branch**?

Satisfiability Problem

Satisfiability

The instances of **Satisfiability** or **SAT**:

$$(x \vee y \vee z)(x \vee \bar{y})(y \vee \bar{z})(z \vee \bar{x})(\bar{x} \vee \bar{y} \vee \bar{z})$$

a Boolean formula in conjunctive normal form (CNF).

It is a collection of clauses (the parentheses),

- each consisting of the disjunction of several literals;
- a literal is either a Boolean variable or the negation of one.

A satisfying truth assignment is an assignment of **false** or **true** to each variable so that every clause contains a literal of **true**.

Given a Boolean formula in conjunctive normal form, either find a satisfying truth assignment or else report that none exists.

2-SAT

Given a set of clauses, where each clause is the disjunction of two literals, looking for an assignment so that all clauses are satisfied.

$$(x_1 \vee x_2) \wedge (\bar{x}_1 \vee x_3) \wedge (x_1 \vee \bar{x}_2) \wedge (x_3 \vee x_4) \wedge (\bar{x}_1 \vee \bar{x}_4)$$

Given an instance I of **2-SAT** with n variables and m clauses, construct a directed graph $G_I = (V, E)$ as follows.

- G_I has $2n$ nodes, one for each **variable** and its negation.
- G_I has $2m$ edges: for each clause $(\alpha \vee \beta)$ of I , G_I has an edge from the negation of α to β , and one from the negation of β to α .

2-SAT

Show that if G_I has a strongly connected component containing both x and \bar{x} for some variable x , then I has no satisfying assignment.

If none of G_I 's strongly connected components contain both a literal and its negation, then the instance I must be satisfiable.

Conclude that there is a linear-time algorithm for solving 2-SAT.

A Little Change of the 2-SAT

3-SAT, SAT?

Search Problems (Decision Problems)

Satisfiability

The instances of **Satisfiability** or **SAT**:

$$(x \vee y \vee z)(x \vee \bar{y})(y \vee \bar{z})(z \vee \bar{x})(\bar{x} \vee \bar{y} \vee \bar{z})$$

a Boolean formula in conjunctive normal form (CNF).

It is a collection of clauses (the parentheses),

- each consisting of the disjunction of several literals;
- a literal is either a Boolean variable or the negation of one.

A satisfying truth assignment is an assignment of **false** or **true** to each variable so that every clause contains a literal of **true**.

Given a Boolean formula in conjunctive normal form, either find a satisfying truth assignment or else report that none exists.

Search Problems

SAT is a typical **search problem**.

We are given an **instance** *I*

- Some input data specifying the problem
- A Boolean formula in conjunctive normal form

we are asked to find a **solution** *S*

- An object that meets a particular specification
- An assignment that satisfies each clause

If no such solution exists, we must say so.

Search Problems

A search problem must have the property that any proposed solution S to an instance I can be quickly checked for correctness.

S must be concise, with length polynomially bounded by that of I .

- This is true for **SAT**, where S is an assignment to the variables.

There is a polynomial-time algorithm that takes as input I and S and decides whether or not S is a solution of I .

- For **SAT**, it is easy to check whether the assignment specified by S satisfies every clause in I .

Search Problems

A **search problem** is specified by an algorithm C that takes two inputs, an **instance** I and a **proposed solution** S , and runs in **time polynomial** in $|I|$.

We say S is a **solution** to I if and only if $C(I, S) = \text{true}$.

Satisfiability Revisit

Researchers over the past 80 years have tried to find efficient ways to solve the **SAT**, but without success.

The fastest algorithms we have are still exponential on their worst-case inputs.

There are two natural variants of **SAT** with good algorithms.

- **2-SAT** can be solved in linear time.
- All clauses contain at most one positive literal, say **Horn formula**, can be found by the greedy algorithm.

Traveling Salesman Problem


The Traveling Salesman Problem

In the **traveling salesman problem(TSP)** we are given n vertices and all $n(n-1)/2$ distances between them, and a **budget b** .

To find a cycle that passes through every vertex exactly once, of total cost b or less - or to report that no such cycle.

A **permutation $\tau(1), \dots, \tau(n)$** of the vertices such that when they are toured in this order, the total distance covered is at most b :

$$d_{\tau(1),\tau(2)} + d_{\tau(2),\tau(3)} + \dots + d_{\tau(n),\tau(1)} \leq b$$



The Traveling Salesman Problem

We have defined the **TSP** as a **search problem**: given an instance, find a tour within the budget (or report that none exists).

But why are we expressing the **TSP** in this way, when in reality it is an **optimization problem**, in which the shortest possible tour is sought?

Search VS. Optimization

Turning an **optimization problem** into a **search problem** does not change its difficulty, because the two versions reduce to one another.

Any algorithm that solves the **optimization** also solves the **search** problem:

- find the **optimum tour** and if it is within **budget**, return it; if not, there is no solution.

Conversely, an algorithm for the **search** problem can also be used to solve the **optimization** problem:

- First suppose that we knew the cost of the optimum tour; then we could find this tour by calling the algorithm for the **search** problem, using the **optimum** cost as the budget.
- We can find the optimum cost by **binary search**.

Search Instead of Optimization

Isn't any optimization problem also a search problem in the sense that we are searching for a solution that has the property of being **optimal**?

The solution to a **search problem** should be easy to recognize, or as we put it earlier, **polynomial-time** checkable.

Given a potential solution to the **TSP**, it is easy to check the properties “**is a tour**” (just check that each vertex is visited exactly once) and “**has total length $\leq b$** .”

But how could one check the property “**is optimal**”?

TSP Revisit

There are no known polynomial-time algorithms for the **TSP**, despite much effort by researchers over nearly a century.

There exists a faster, yet still exponential, **dynamic programming** algorithm.

The **Minimum spanning tree (MST)** problem, for which we do have **efficient algorithms**, provides a stark contrast here.


The **TSP** can be thought of as a tough cousin of the **MST** problem, in which the tree is not allowed to **branch** and is therefore a **path**.

This extra restriction on the structure of the tree results in a much **harder problem**.

Euler and Rudrata

Euler path:

Given a graph, find a path that contains each edge exactly once.




Euler Path

The answer is **yes** if and only if

- ① the graph is **connected** and
- ② every vertex, with the possible exception of two vertices (the start and final vertices of the walk), has **even degree**.

There is a **polynomial time** algorithm for **Euler path**.

Rudrata Cycle




Rudrata Cycle

Rudrata Cycle:

Given a graph, find a cycle that visits each vertex exactly once.

In the literature this problem is known as the **Hamilton cycle problem**.




Cuts and Bisections

Minimum Cut

A **cut** is a set of edges whose removal leaves a graph disconnected.

Minimum cut: given a graph and a budget b , find a cut with at most b edges.



Minimum Cut

This problem can be solved in polynomial time by $n - 1$ max-flow computations:

- give each edge a capacity of 1,
- and find the maximum flow between some fixed node and every single other node.

The smallest such flow will correspond (via the max-flow min-cut theorem) to the smallest cut.

Balanced Cut

In many graphs, the smallest cut leaves just a **singleton** vertex on one side - it consists of all edges **adjacent** to this vertex.

Far more interesting are small cuts that partition the vertices of the graph into nearly equal-sized sets.

Balanced cut: Given a graph with n vertices and a budget b , partition the vertices into two sets S and T such that $|S|, |T| \geq n/3$ and such that there are at most b edges between S and T .

Integer Linear Programming

Linear Programming

In a **linear programming problem** we are given a set of **variables**, and to assign **real values** to them so as to

- ① satisfy a set of **linear equations** and/or **linear inequalities** involving these variables, and
- ② maximize or minimize a given **linear objective function**.

Linear Programming

$$\max x_1 + 6x_2 + 13x_3$$

$$x_1 \leq 200$$

$$x_2 \leq 300$$

$$x_1 + x_2 + x_3 \leq 400$$

$$x_2 + 3x_3 \leq 600$$

$$x_1, x_2, x_3 \geq 0$$

Integer Linear Programming

Integer linear programming (ILP): We are given a set of linear inequalities $A\mathbf{x} \leq \mathbf{b}$, where

- A is an $m \times n$ matrix and
- \mathbf{b} is an m -vector;
- an **objective function** specified by an n -vector \mathbf{c} ;
- a goal g .

We want to find a nonnegative integer n -vector \mathbf{x} such that $A\mathbf{x} \leq \mathbf{b}$ and $\mathbf{c} \cdot \mathbf{x} \geq g$.

Integer Linear Programming

$$\max 2x_1 + 5x_2$$

$$2x_1 - x_2 \leq 4$$

$$x_1 + 2x_2 \leq 9$$

$$-x_1 + x_2 \leq 3$$

$$x_1, x_2 \geq 0$$

$$2x_1 + 5x_2 \leq g$$

$$2x_1 - x_2 \leq 4$$

$$x_1 + 2x_2 \leq 9$$

$$-x_1 + x_2 \leq 3$$

$$x_1, x_2 \geq 0$$

But there is a redundancy here:

- the last constraint $c \cdot \mathbf{x} \geq g$ is itself a **linear inequality** and
- can be **absorbed** into $A\mathbf{x} \leq b$.

Integer Linear Programming

So, we define **ILP** to be following **search problem**:

Given A and b , find a nonnegative integer vector x satisfying the inequalities $Ax \leq b$.

3D-Matching

Bipartite Matching

BOYS

GIRLS

Al

Alice

Bob


Beatrice

Chet


Carol

Dan

Danielle



Bipartite Matching




Three-Dimensional Matching

3D matching: There are n boys, n girls, and n pets. The compatibilities are specified by a set of **triples**, each containing a boy, a girl, and a pet.

A triple (b, g, p) means that boy b , girl g , and pet p get along well together.

To find n disjoint triples and thereby create n harmonious households.




Independent Set

Independent Set, Vertex Cover, and Clique

Independent set: Given a graph and an integer g , find g vertices, no two of which have an edge between them.

Vertex cover: Given a graph and an integer b , find b vertices cover (touch) every edge.

Clique: Given a graph and an integer g , find g vertices such that all possible edges between them are present.



Longest Path

Longest Path

Longest path: Given a graph G with nonnegative edge weights and two distinguished vertices s and t , along with a goal g .

To find a path from s to t with total weight at least g .

To avoid trivial solutions we require that the path be simple, containing no repeated vertices.

Knapsack

Knapsack

Knapsack: Given integer weights w_1, \dots, w_n and integer values v_1, \dots, v_n for n items. We are also given a weight capacity W and a goal g .

Seek a set of items whose total weight is **at most** W and whose total value is **at least** g .

The problem is solvable in time $O(nW)$ by **dynamic programming**.

Knapsack

Is there a polynomial algorithm for Knapsack? Nobody knows of one.

A variant of the Knapsack problem is that the **unary** integers.

- by writing **IIIIIIIIII** for **12**.
- It defines a legitimate problem, which we could call **Unary knapsack**.
- It has a **polynomial algorithm**.

A different variation:

- Suppose now that each item's value is equal to its weight, the goal **g** is the same as the capacity **W**.
- This special case is tantamount to finding a subset of a given set of integers that adds up to exactly **W**.
- **Q**: Could it be polynomial?

Subset Sum

Subset sum: Find a subset of a given set of integers that adds up to exactly W .

Referred Materials

- Content of this lecture comes from Section 8.1, 5.1, and 3.4 in [DPV07].
- Suggest to read Chapter 34 of [CLRS09].