

Design and Analysis of Algorithms (VIII)

An Introduction to Approximation Algorithms

Guoqiang Li

School of Software, Shanghai Jiao Tong University

NP-Hard and Optimization Problems

Combinatorial optimization is a topic that consists of finding an optimal object from a finite set of objects.

Most natural optimization problems, including those arising in application areas, are **NP-hard**. Exhaustive search is not feasible.

Under the widely believed conjecture that $P \neq NP$, their exact solution is prohibitively time consuming.


Approximability of these problems becomes a compelling subject of scientific inquiry in computer science and mathematics.

Vertex cover

Vertex cover

Given an undirected graph $G = (V, E)$, and a cost function on vertices $c : V \rightarrow \mathbb{Q}^+$, find a minimum cost **vertex cover**, i.e., a set $V' \subseteq V$ such that every edge has at least one endpoint incident at V' .

The special case, in which all vertices are of unit cost, will be called the **cardinality vertex cover** problem.



NP-Optimization Problem

An **NP-optimization problem** Π is either a minimization or a maximization problem.

Each valid **instance** I of Π comes with a nonempty set of feasible solutions, each of which is assigned a nonnegative rational number called its **objective function value**.

There exist **polynomial time algorithms** for determining validity, feasibility, and the objective function value.

A feasible solution that achieves the optimal objective function value is called an **optimal solution**.

NP-Optimization Problem

$OPT_{\Pi}(I)$ denotes the **objective function value** of an optimal solution to **instance I** . OPT is used when there is no ambiguity.

An **approximation algorithm**, \mathcal{A} , for Π is in **polynomial time**. A **feasible solution** of objective function value is “close” to the optimal.

By “close” we mean within a **guaranteed factor** of the optimal.

A Dilemma

To establish the **approximation guarantee**, the cost of the solution produced by the algorithm needs to compare with an **optimal solution**.

For such problems, not only is it NP-hard to find an **optimal solution**, but it is also NP-hard to compute the **cost** of an optimal solution.

In fact, computing the cost of an optimal solution is precisely the **difficult core** of such problems.

How do we establish the **approximation guarantee**? The answer provides a key step in the design of **approximation algorithms**.



Cardinality Vertex Cover

Matching

Given a graph $G = (V, E)$, a subset of the edges $M \subseteq E$ is said to be a **matching** if no two edges of M share an endpoint.

A matching of **maximum cardinality** in G is called a **maximum matching**.

A matching that is **maximal under inclusion** is called a **maximal matching**.



Matching

Given a graph $G = (V, E)$, a subset of the edges $M \subseteq E$ is said to be a **matching** if no two edges of M share an endpoint.

A matching of **maximum cardinality** in G is called a maximum matching.



A matching that is maximal under inclusion is called a maximal matching.

A **maximal matching** can clearly be computed in **polynomial time** by simply **greedily** picking edges and removing endpoints of picked edges. More sophisticated means lead to polynomial time algorithms for finding a **maximum matching** as well.

Approximation for Cardinality VC

Algorithm

Find a **maximal matching** in G and output the set of **matched vertices**.



Approximation Factor



The **Algorithm** is a **factor 2** approximation algorithm for the **cardinality vertex cover** problem.

Proof.

- No edge can be left uncovered by the set of vertices picked.
- Let M be the matching picked. As argued above,

$$|M| \leq OPT$$

- The approximation factor is at most $2 \cdot OPT$.



Lower Bounding OPT

The approximation algorithm for vertex cover was very much related to, and followed naturally from, the **lower bounding scheme**. This is in fact typical in the design of approximation algorithms.

Can the Guarantee be Improved?


Can the **approximation guarantee** of **Algorithm** be improved by a better analysis?

Can an approximation algorithm with a **better guarantee** be designed using the **lower bounding scheme** of **Algorithm**?

Is there some other lower bounding method that can lead to an improved approximation guarantee for **Vertex cover**?

A Better Analysis?

Consider the infinite family of instances given by the **complete bipartite graphs** $K_{n,n}$.



When run on $K_{n,n}$, **Algorithm** will pick all $2n$ vertices, whereas picking one side of the bipartition gives a **cover** of size n .

Tight Example

$K_{n,n}$ shows that the analysis is **tight**, by giving an infinite family of instances in which the solution is **twice** the optimal.

An **infinite family** of instances showing that the analysis of an approximation algorithm is tight, is referred to as a **tight example**.

Tight examples for an approximation algorithm give critical insight into the functioning of the algorithm.

They have often led to ideas for obtaining algorithms with **improved guarantees**.

A Better Guarantee?

The lower bound, of size of a **maximal matching**, is half the size of an optimal **vertex cover** for the following infinite family of instances. Consider the **complete graph** K_n , where n is odd. The size of any maximal matching is $(n - 1)/2$, whereas the size of an optimal cover is $n - 1$.

A Better Algorithm?

Still Open!

Set Cover

Set cover

Set cover

Given a **universe** U of n elements, a **collection** of subsets of U , $\mathcal{S} = \{S_1, \dots, S_k\}$, and a **cost function** $c : \mathcal{S} \rightarrow \mathbb{Q}^+$, find a **minimum cost sub-collection** of \mathcal{S} that covers all elements of U .

Some Remarks

Define the **frequency** f of an element to be the number of sets it is in.

The various **approximation algorithms** for **set cover** achieve one of two factors: $O(\log n)$ or f .

Clearly, neither **dominates** the other in all instances.

The special case of **set cover** with $f = 2$ is essentially the **vertex cover** problem, for which we gave a factor 2 approximation algorithm.

Cardinality Set Cover

The Problem

A county is in its early stages of planning and is deciding where to put schools.

There are only two constraints:


- each school should be **in a town**,
- and no one should have to travel more than **30** miles to reach one of them.

Q: What is the minimum number of schools needed?

(a)



(b)



The Problem

This is a typical (cardinality) set cover problem.


- For each town x , let S_x be the set of towns within 30 miles of it.
- A school at x will essentially “cover” these other towns.
- The question is then, how many sets S_x must be picked in order to cover all the towns in the county?

Set Cover Problem

Set Cover

- **Input:** A set of elements B , sets $S_1, \dots, S_m \subseteq B$
- **Output:** A selection of the S_i whose union is B .
- **Cost:** Number of sets picked.

The Example



Performance Ratio

Lemma

Suppose B contains n elements and that the optimal cover consists of OPT sets. Then the greedy algorithm will use at most $\ln n \cdot OPT$ sets.

Proof:

Let n_t be the number of elements still not covered after t iterations of the greedy algorithm (so $n_0 = n$).

Since these remaining elements are covered by the optimal OPT sets, there must be some set with at least n_t / OPT of them.

Therefore, the greedy strategy will ensure that

$$n_{t+1} \leq n_t - \frac{n_t}{OPT} = n_t \left(1 - \frac{1}{OPT}\right)$$

which by repeated application implies

$$n_t \leq n_0 \left(1 - \frac{1}{OPT}\right)^t$$

Performance Ratio

A more convenient bound can be obtained from the useful inequality

$$1 - x \leq e^{-x} \text{ for all } x$$

with equality if and only if $x = 0$,

Thus

$$n_t \leq n_0 \left(1 - \frac{1}{OPT}\right)^t < n_0 \left(e^{-\frac{1}{OPT}}\right)^t = ne^{-\frac{t}{OPT}}$$

At $t = \ln n \cdot OPT$, therefore, n_t is strictly less than $ne^{-\ln n} = 1$, which means no elements remain to be covered.

Set Cover

The Greedy Algorithm

The Greedy Strategies

Iteratively pick the most **cost-effective set** and remove the covered elements, until all elements are covered.

Let C be the set of elements already covered at the beginning of an iteration.

During this iteration, define the **cost-effectiveness** of a set S to be the **average cost** at which it covers new elements, i.e., $\frac{c(S)}{|S - C|}$.

Define the **price** of an element to be the average cost at which it is covered.

When a set S is picked, we can think of its **cost** being distributed equally among the new elements covered, to set their prices.

The Greedy Algorithm

- ① $C \leftarrow \emptyset$.
- ② While $C \neq U$ do
 - Find the most cost-effective set in the current iteration, say S .
 - Let $\alpha = \frac{c(S)}{|S-C|}$, i.e., the cost-effectiveness of S .
 - Pick S , and for each $e \in S - C$, set $price(e) = \alpha$.
 - $C \leftarrow C \cup S$.
- ③ Output the picked sets.

The Lemma

Lemma

Number the elements of U in the order in which they were covered by the algorithm, resolving ties arbitrarily. Let e_1, \dots, e_n be this numbering.

For each $k \in \{1, \dots, n\}$, $price(e_k) \leq OPT/(n - k + 1)$.

The Lemma

Proof.

In any iteration, the **leftover sets** of the optimal solution can cover the **remaining elements** at a cost of at most OPT .

Therefore, among these sets, there must be one having **cost-effectiveness** of at most $OPT/|\overline{C}|$.

$$price(e_k) \leq \frac{OPT}{|\overline{C}|}$$

In the iteration in which element e_k was covered, \overline{C} contained **at least** $n - k + 1$ elements.

$$price(e_k) \leq \frac{OPT}{|\overline{C}|} \leq \frac{OPT}{n - k + 1}$$

The Approximation Factor


Theorem

The greedy algorithm is an H_n factor approximation algorithm for the minimum set cover problem, where

$$H_n = 1 + \frac{1}{2} + \dots + \frac{1}{n}$$

A Tight Example

The following is a tight example



When run on this instance the greedy algorithm outputs the cover consisting of the n singleton sets, since in each iteration some singleton is the most cost-effective set. Thus, the algorithm outputs a cover of cost

$$= \frac{1}{n} + \frac{1}{n-1} + \dots + 1 = H_n$$

On the other hand, the **optimal cover** has a **cost** of $1 + \epsilon$.

Layering

Layering

The algorithm design technique of layering is also best introduced via set cover. However, that this is not a very widely applicable technique.

We will give a factor 2 approximation algorithm for **vertex cover**, assuming **arbitrary weights**.

The idea in layering is to decompose the **given weight function** on vertices into convenient functions, called **degree-weighted**, on a nested sequence of subgraphs of G .

The Lemma

Let $\omega : V \rightarrow \mathbb{Q}^+$ be the function assigning **weights** to the vertices of the given graph $G = (V, E)$.

A function assigning vertex weights is **degree-weighted** if there is a constant $c > 0$ such that the weight of each vertex $v \in V$ is $c \cdot \deg(v)$.

Lemma

In **Vertex Cover**, let $\omega : V \rightarrow \mathbb{Q}^+$ be a **degree-weighted** function. Then $\omega(V) \leq 2 \cdot OPT$.

The Lemma

Lemma

In **Vertex Cover**, let $\omega : V \rightarrow \mathbb{Q}^+$ be a **degree-weighted** function. Then $\omega(V) \leq 2 \cdot OPT$.

Proof.

Let c be the constant such that $\omega(v) = c \cdot \deg(v)$, and let U be an **optimal vertex cover** in G .

Since U covers all the edges, $\sum_{v \in U} \deg(v) \geq |E|$.

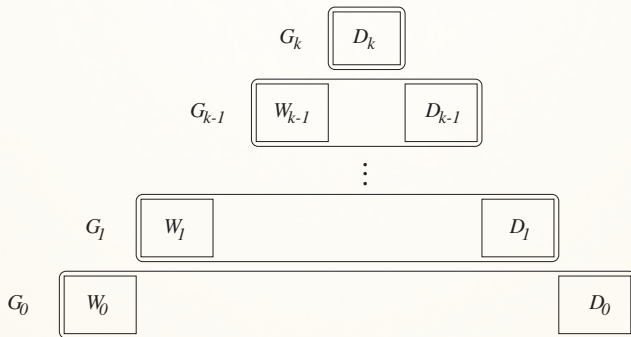
Therefore, $\omega(U) \geq c|E|$. Since $\sum_{v \in V} \deg(v) = 2|E|$, $\omega(V) = 2c|E|$.

The lemma follows.


The Layer Algorithm

- 1 $G_0 = G, C = \emptyset, i = 0.$
- 2 Remove **degree zero** vertices from G_i , say this set is D_i .
- 3 Compute $c = \min\{w(v)/\deg(v)\}$ for all $v \in G_i$.
- 4 Let $t_i(v) = c \cdot \deg(v)$ and $w(v) = w(v) - t_i(v)$ for all $v \in G_i$.
- 5 Let $W_i = \{v \in G_i \mid w(v) = 0\}, C = C \cup W_i$.
- 6 Let G_{i+1} be the graph induced by $V_i - (D_i \cup W_i)$. Increase i by 1 and goto step **2** until G_i is empty graph.


The Layer Algorithm




An Example




An Example




An Example




An Example



An Example



An Example



Analysis

Theorem

The **layer algorithm** achieves an approximation guarantee of factor 2 for the vertex cover problem, assuming **arbitrary vertex weights**.

Proof.

- Firstly, C is a **vertex cover** for G .
- Otherwise, there must be some $(u, v) \in E$ with $u \in D_i$ and $v \in D_j$.
- Assume $i \leq j$, then (u, v) is in G_i contradicting the fact that u is of **degree zero**.

Analysis (cont'd)

Then we show $\omega(c) \leq 2 \cdot \text{OPT}$. Let C^* be an **optimal vertex cover**.

- For $v \in C$, if $v \in W_j$

-

$$\omega(v) = \sum_{i \leq j} t_i(v)$$

- For $v \in V - C$, if $v \in D_j$, then

-

$$\omega(v) \geq \sum_{i < j} t_i(v)$$


Analysis (cont'd)

- In each layer i , $C^* \cap G_i$ is a **vertex cover** for G_i .
- Thus by previous lemma, $t_i(C \cap G_i) \leq 2 \cdot t_i(C^* \cap G_i)$.
- Therefore,

$$\omega(C) = \sum_{i=0}^{k-1} t_i(C \cap G_i) \leq 2 \sum_{i=0}^{k-1} t_i(C^* \cap G_i) \leq 2 \cdot \omega(C^*)$$

A Tight Example

A tight example is provided by the family of **complete bipartite graphs**, $K_{n,n}$, with all vertices of **unit weight**. The **layering algorithm** will pick all $2n$ vertices of $K_{n,n}$ in the cover, whereas the **optimal cover** picks only one side of the bipartition.



Referred Materials

Content of this lecture comes from Chapter 1 and 2 in [Vaz04].

Suggest to read the rest part of Chapter 1 and 2 in [Vaz04].