References

References

1. M. Schwarzkopf. **The evolution of cluster scheduler architectures.** http://www.firmament.io/blog/scheduler-architectures.html.
2. A. Wang. **Mesos, Omega, Borg: A Survey.** https://www.umbrant.com/2015/05/27/mesos-omega-borg-a-survey/.
3. R. Olaniyan and M. Maheswaran. **Recent Developments in Resource Management in Cloud Computing and Large Computing Clusters.** In *Research Advances in Cloud Computing*, 11(4):237-261, 2017.
4. B. Hindman, A. Konwinski, M. Zaharia, A. Ghodsi, A. D. Joseph, R. H. Katz, S. Shenker, and I. Stoica. **Mesos: a platform for fine-grained resource sharing in the data center.** In *NSDI '11*, 2011.
5. J. Dean and S. Ghemawat. **MapReduce: Simplified Data Processing on Large Clusters.** In *Proc. OSDI '04*, 2004.
6. V. K. Vavilapalli, S. Seth, B. Saha, C. Curino, O. O'Malley, S. Radia, B. Reed, E. Baldeschwieler, A. C. Murthy, C. Douglas, S. Agarwal, M. Konar, R. Evans, T. Graves, J. Lowe, and H. Shah. **Apache Hadoop YARN: Yet Another Resource Negotiator.** In *SoCC '13*, 2013.
7. M. Schwarzkopf, A. Konwinski, M. Abd-El-Malek, and J. Wilkes. **Omega: flexible, scalable schedulers for large compute clusters.** In *EuroSys '13*, 2013, pp. 351–364.
8. E. Boutin, J. Ekanayake, W. Lin, B. Shi, J. Zhou, Z. Qian, M. Wu, and L. Zhou. **Apollo: Scalable and Coordinated Scheduling for Cloud-Scale Computing.** In *Proc. OSDI '04*, 2004, pp. 285–300.
9. K. Ousterhout, P. Wendell, M. Zaharia, and I. Stoica. **Sparrow: Distributed, Low Latency Scheduling.** In *SOSP '13*, 2013, pp. 69–84.
10. A. Verma, L. Pedrosa, M. Korupolu, D. Oppenheimer, E. Tune, and J. Wilkes. **Large-scale cluster management at Google with Borg.** In *EuroSys '15*, 2015, pp. 1–17.

References

11. M. Mitzenmacher. **The power of two choices in randomized load balancing.** In *IEEE Transactions on Parallel and Distributed Systems*, vol. 12, no. 10, pp. 1094–1104, 2001.
12. G. Park. **A Generalization of Multiple Choice Balls-into-Bins.** In *PODC '11*, 2011, pp. 297–298.
13. P. Delgado, F. Dinu, A.-M. Kermarrec, and W. Zwaenepoel. **Hawk: Hybrid Datacenter Scheduling.** In *USENIX ATC '15*, 2015.
14. K. Karanasos, S. Rao, C. Curino, C. Douglas, K. Chaliparambil, G. M. Fumarola, S. Heddaya, R. Ramakrishnan, and S. Sakalanaga, **Mercury: Hybrid Centralized and Distributed Scheduling in Large Shared Clusters.** In *USENIX ATC '15*, 2015.
15. **Running Spark on Kubernetes.** https://spark.apache.org/docs/latest/running-on-kubernetes.html.
16. **Apache Spark 2.3 with Native Kubernetes Support.** https://kubernetes.io/blog/2018/03/apache-spark-23-with-native-kubernetes/.